

Machine Learning CS 529 Project 1 Report

Esteban Guillen

High-level Description of code

I implemented the ID3 algorithm in Python (Figure 1). All code is contained in a file called *id3.py*. The code follows the algorithm (https://en.wikipedia.org/wiki/ID3_algorithm) closely and adds the ability to switch between entropy and misclassification error for the impurity measure. Split stopping using the Chi-squared test was also added.

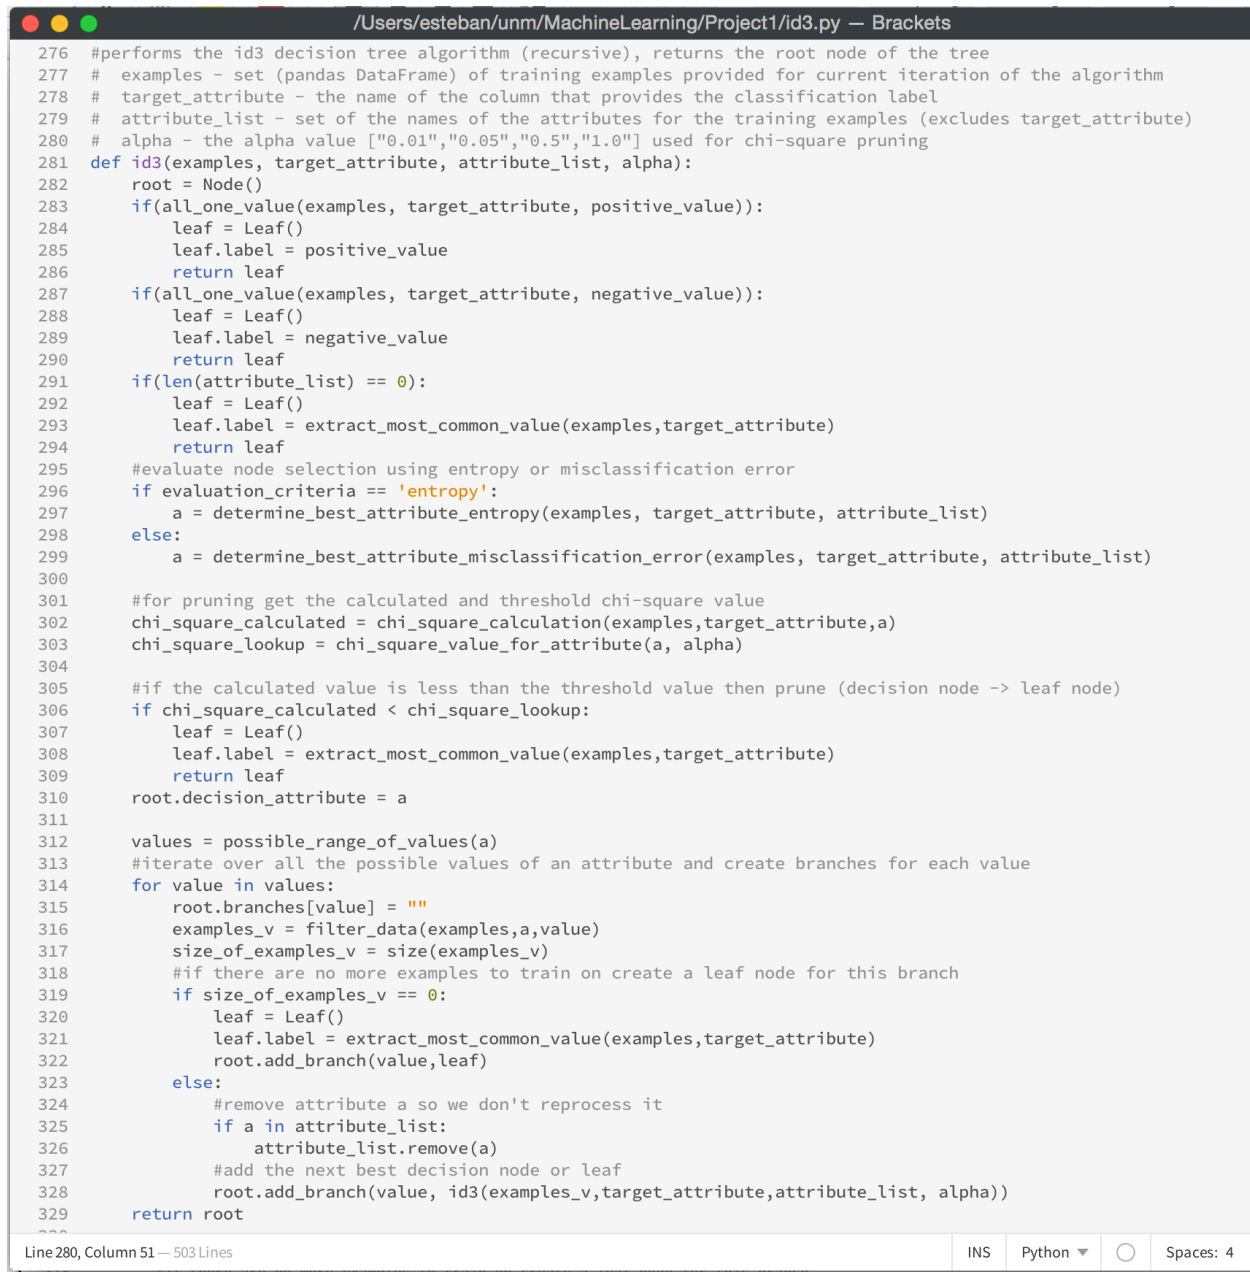
The image shows a screenshot of a code editor window titled "/Users/esteban/unm/MachineLearning/Project1/id3.py — Brackets". The editor displays Python code for the ID3 decision tree algorithm. The code is line-numbered from 276 to 329. It includes comments explaining the parameters and the logic of the algorithm. The code defines a function 'id3' that takes 'examples', 'target_attribute', 'attribute_list', and 'alpha' as inputs. It recursively builds a decision tree by selecting the best attribute to split on based on entropy or misclassification error, and then pruning the tree using a chi-square test. The code is well-commented and uses standard Python syntax for loops, conditionals, and function calls. The editor interface at the bottom shows 'Line 280, Column 51 — 503 Lines', 'INS', 'Python', and 'Spaces: 4'.

Figure 1. ID3 Algorithm

I used pandas DataFrame data structures for holding the training, testing and validation data (Figure 2 shows training and testing data being loaded). The DataFrame made it easy to index into the data and create subsets based on attribute values.

```

/Users/esteban/unm/MachineLearning/Project1/id3.py — Brackets
7  names_of_attributes = names=["label","cap-shape","cap-surface","cap-color","bruises","oder",
8                                "gill-attachment","gill-spacing","gill-size","gill-color","stalk-shape",
9                                "stalk-root","stalk-surface-above-ring","stalk-surface-below-ring",
10                               "stalk-color-above-ring","stalk-color-below-ring","veil-type",
11                               "veil-color","ring-number","ring-type","spore-print-color",
12                               "population","habitat"]
13
14  #read training data into pandas DataFrame
15  training_data = pd.read_table("data/training.txt", sep=",", names=names_of_attributes)
16
17  #read testing data into pandas DataFrame
18  testing_data = pd.read_table("data/testing.txt", sep=",", names=names_of_attributes)
19
Line 280, Column 51 — 503 Lines    INS    Python ▾    ○    Spaces: 4
```

Figure 2. Using pandas DataFrame

I used two types of objects to represent my decision tree (Figure 3). **Node** which is a decision node that consists of a string representing the decision attribute and a dictionary to manage the branches and child nodes. **Leaf** which simply consists of a string that represents the classification label (e or p).

```

/Users/esteban/unm/MachineLearning/Project1/id3.py — Brackets
21  #represents a decision node in the tree
22  class Node:
23      #identifies which attribute this node tests on
24      decision_attribute = ""
25
26      #initialize node
27      def __init__(self):
28          self.decision_attribute = ""
29          self.branches = {}
30
31      #adds a child Node or Leaf
32      def add_branch(self,attribute_value,node):
33          self.branches[attribute_value] = node
34
35      #returns all branches
36      def get_branches(self):
37          return self.branches
38
39  #represents a leaf node in the tree
40  class Leaf:
41      #label representing the classification value (e or p)
42      label = ""
43
Line 280, Column 51 — 503 Lines    INS    Python ▾    ○    Spaces: 4
```

Figure 3. Objects used to represent the decision tree

The calculations for entropy and misclassification error can be seen below (Figure 4). The calculations were very straight forward (entropy used the log function from the math library).

```
179 #calculates the entropy on a subset (examples_v) of data
180 def calculate_entropy(examples_v,target_attribute):
181     entropy = 0.0
182     total_size = size(examples_v)
183     #if there is no data just return a zero value (does not contribute to calculation)
184     if total_size == 0:
185         return 0.0
186     positive_filtered_examples_v = filter_data(examples_v, target_attribute, positive_value)
187     negative_filtered_examples_v = filter_data(examples_v, target_attribute, negative_value)
188
189     num_positive = size(positive_filtered_examples_v)
190     calc_positive = 0.0
191     if num_positive != 0:
192         calc_positive = -(num_positive/total_size)*math.log((num_positive/total_size),2)
193
194     num_negative = size(negative_filtered_examples_v)
195     calc_negative = 0.0
196     if num_negative != 0:
197         calc_negative = - (num_negative/total_size)*math.log((num_negative/total_size),2)
198
199     #returns -p(+)log p(+) - p(-)log p(-)
200     return calc_positive + calc_negative
201
202
203 #calculates the misclassification error on a subset (examples_v) of data
204 def calculate_misclassification_error(examples_v, target_attribute):
205     error = 0.0
206     total_size = size(examples_v)
207     #if there is no data just return a zero value (does not contribute to calculation)
208     if total_size == 0:
209         return 0.0
210
211     positive_filtered_examples_v = filter_data(examples_v,target_attribute,positive_value)
212     negative_filtered_examples_v = filter_data(examples_v,target_attribute,negative_value)
213
214     num_positive = size(positive_filtered_examples_v)
215     p_positive = (num_positive/total_size)
216
217     num_negative = size(negative_filtered_examples_v)
218     p_negative = (num_negative/total_size)
219
220     error = 1.0 - max([p_positive,p_negative])
221     #returns 1 - max(probability of positive,probability of negative)
222     return error
```

Line 217, Column 54 — 503 Lines INS Python ▼ ○ Spaces: 4

Figure 4. Entropy and misclassification calculations

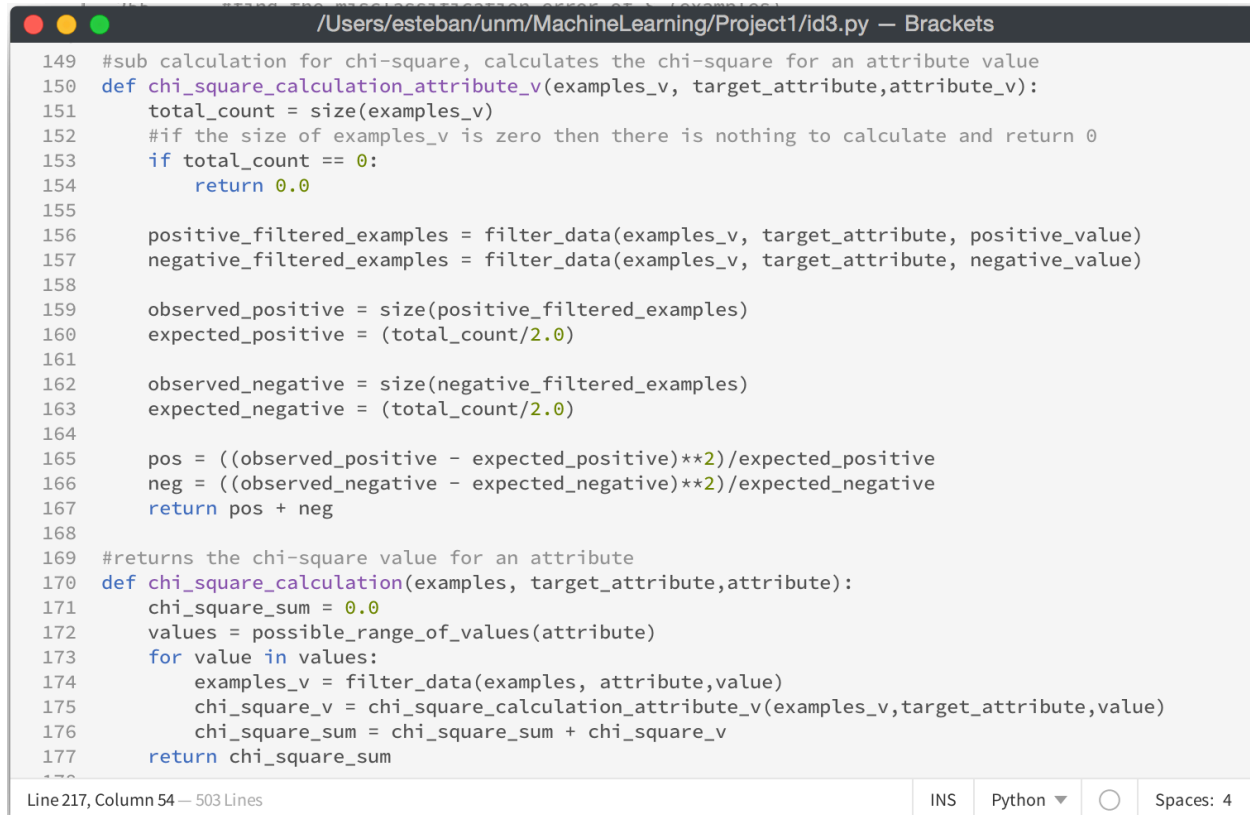
Gain was used to determine the “best attribute” for each recursive call of the ID3 algorithm. The implementations for entropy and misclassification can be seen below (Figure 5).

A screenshot of a code editor window titled "/Users/esteban/unm/MachineLearning/Project1/id3.py — Brackets". The editor displays two Python functions. The first function, `determine_best_attribute_entropy`, starts at line 224 and ends at line 250. It calculates the entropy of a set of examples and then iterates through attributes to find the one that maximizes information gain. The second function, `determine_best_attribute_misclassification_error`, starts at line 251 and ends at line 274. It calculates the misclassification error of a set of examples and then iterates through attributes to find the one that maximizes the gain (minimizes error). The code is well-commented and uses standard Python syntax for loops, conditionals, and function definitions. The editor interface includes a status bar at the bottom showing "Line 217, Column 54 — 503 Lines", "INS", "Python", and "Spaces: 4".

```
224 #returns the attribute with the best informatin gain using entropy as the impurity measure
225 def determine_best_attribute_entropy(examples, target_attribute, attributes):
226     best_attribute = ""
227     max_information_gain = 0.0
228
229     #find the entropy of S (examples)
230     entropy_s = calculate_entropy(examples,target_attribute)
231     size_of_s = size(examples)
232
233
234     #loop through attributes and calculate the gain
235     for attribute in attributes:
236         values = possible_range_of_values(attribute)
237         information_gain = entropy_s
238         for value in values:
239             examples_v = filter_data(examples,attribute,value)
240             entropy_v = calculate_entropy(examples_v,target_attribute)
241             size_of_examples_v = size(examples_v)
242             weight = (size_of_examples_v/size_of_s)
243             information_gain = information_gain - weight*entropy_v
244             #if we have a new candidate for the best attribute store the gain and attribute name
245             if information_gain > max_information_gain:
246                 max_information_gain = information_gain
247                 best_attribute = attribute
248
249     return best_attribute
250
251 #returns the attribute with the best gain using misclassification error as the impurity measure
252 def determine_best_attribute_misclassification_error(examples, target_attribute, attributes):
253     best_attribute = ""
254     max_gain = 0.0
255     #find the misclassification error of S (examples)
256     error_s = calculate_misclassification_error(examples,target_attribute)
257     size_of_s = size(examples)
258
259     #loop through attributes and calculate the gain
260     for attribute in attributes:
261         values = possible_range_of_values(attribute)
262         gain = error_s
263         for value in values:
264             examples_v = filter_data(examples,attribute,value)
265             error_v = calculate_misclassification_error(examples_v,target_attribute)
266             size_of_examples_v = size(examples_v)
267             weight = (size_of_examples_v/size_of_s)
268             gain = gain - weight*error_v
269             #if we have a new candidate for the best attribute store the gain and attribute name
270             if gain > max_gain:
271                 max_gain = gain
272                 best_attribute = attribute
273
274     return best_attribute
---
```

Figure 5. Gain calculations for entropy and misclassification

The Chi-square test was used for split stopping and the code for calculating the Chi-square value is shown below (Figure 6). The calculation was broken up into 2 functions, the first calculates the Chi-square for a specific value of an attribute and the second sums all those calculations (over all possible values of an attribute).

A screenshot of a code editor window titled "/Users/esteban/unm/MachineLearning/Project1/id3.py — Brackets". The editor displays Python code for calculating the Chi-square value. The code is organized into two functions. The first function, `chi_square_calculation_attribute_v`, calculates the Chi-square for a specific attribute value by filtering examples, calculating observed and expected counts, and then computing the Chi-square contribution for positive and negative values. The second function, `chi_square_calculation`, iterates over all possible values of an attribute and sums the Chi-square values calculated by the first function. The code is well-commented and includes a return statement for the total Chi-square sum.

```
149 #sub calculation for chi-square, calculates the chi-square for an attribute value
150 def chi_square_calculation_attribute_v(examples_v, target_attribute, attribute_v):
151     total_count = size(examples_v)
152     #if the size of examples_v is zero then there is nothing to calculate and return 0
153     if total_count == 0:
154         return 0.0
155
156     positive_filtered_examples = filter_data(examples_v, target_attribute, positive_value)
157     negative_filtered_examples = filter_data(examples_v, target_attribute, negative_value)
158
159     observed_positive = size(positive_filtered_examples)
160     expected_positive = (total_count/2.0)
161
162     observed_negative = size(negative_filtered_examples)
163     expected_negative = (total_count/2.0)
164
165     pos = ((observed_positive - expected_positive)**2)/expected_positive
166     neg = ((observed_negative - expected_negative)**2)/expected_negative
167     return pos + neg
168
169 #returns the chi-square value for an attribute
170 def chi_square_calculation(examples, target_attribute, attribute):
171     chi_square_sum = 0.0
172     values = possible_range_of_values(attribute)
173     for value in values:
174         examples_v = filter_data(examples, attribute, value)
175         chi_square_v = chi_square_calculation_attribute_v(examples_v, target_attribute, value)
176         chi_square_sum = chi_square_sum + chi_square_v
177     return chi_square_sum
```

Line 217, Column 54 — 503 Lines INS Python ▼ ○ Spaces: 4

Figure 6. Chi-square calculation code

I implemented a number of helper functions to make the code more readable and maintainable. Check the full source code (*id3.py*) for more details.

Accuracy Results

All of my decision trees (entropy and misclassification error for 99, 95, 50, and 0 CL) produced 100% accuracy results (Figure 7). None of my trees got pruned using the Chi-squared test, all of my calculated values were less than the threshold values. The trees that were produced were small (4 levels for entropy and 5 levels for misclassification error, with most decision nodes having only one branch pointing to another decision node). The training data provided must have produced a near optimal tree, and no amount of pruning would improve performance (can't get better than 100%). Most of the training data was classified at the root node and 8 of the 9 branches pointed to Leaf nodes.

```
Project1 — bash — 90x34
python  ...  bash  bash  bash  +

Estebans-MacBook-Pro:Project1 esteban$ python3 id3.py

Accuracy (using entropy and confidence level of 0):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using entropy and confidence level of 99):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using entropy and confidence level of 95):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using entropy and confidence level of 50):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0

Accuracy (using misclassification and confidence level of 0):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using misclassification and confidence level of 99):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using misclassification and confidence level of 95):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0
Accuracy (using misclassification and confidence level of 50):  100.0 %
  Number correctly classified:  2031
  Number incorrectly classified:  0

best tree:  id3_entropy_confidence_level_0

classification predictions on validation data written to: validation-best-accuracy.txt

Estebans-MacBook-Pro:Project1 esteban$
```

Figure 7. Classification results on the training.txt data

I used the entropy with CL 0% decision tree to classify the validation data. The 100% classification results can be seen below (Figure 8)

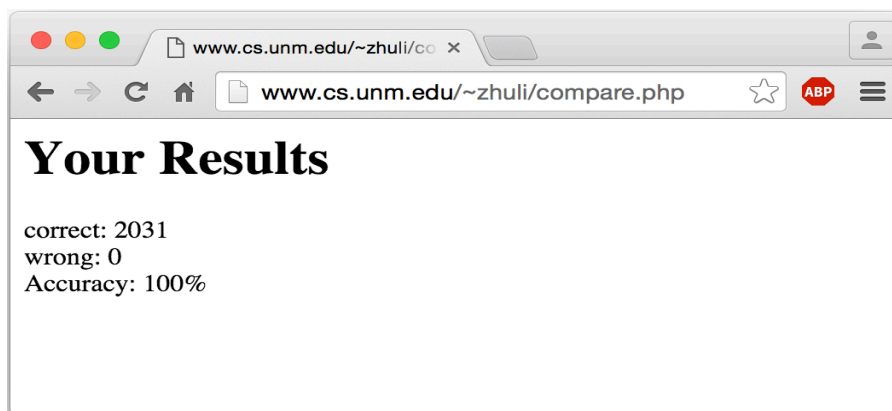


Figure 8. Accuracy for the validation.txt data

Folklore Rules

For proving or disproving the folklore rules I look at the data and/or my decision tree to come to a conclusion. I concluded that a folklore rule was true if the data or my decision tree backed-up the claim with a high probability (most of the data supported the claim). The included HTML export of my IPython Notebook shows how I queried the data.

Rule 1: Poisonous mushrooms are brightly colored:

False: I found there to be a fairly even split between edible and poisonous mushrooms for brightly colored (cap-color) mushrooms.

Rule 2: Poisonous mushrooms taste/smell bad:

True: The data provided strongly supports this (for smell bad). Almost all mushrooms with a bad odor were poisonous. My decision tree also supports this rule.

Rule 3: Poisonous mushrooms have a pointed or umbrella shaped cap:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having a pointed or umbrella shaped cap.

Rule 4: Edible mushroom have flat rounded shaped cap:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having a flat shaped cap.

Rule 5: Poisonous mushrooms have warts or scales on the cap:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having warts or scales on the cap.

Rule 6: Poisonous mushrooms have a bulbous cup or sac around the base:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having a bulbous cup or sac around the base.

Rule 7: Poisonous mushroom have a ring around the stem:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having a ring around the stem.

Rule 8: Poisonous mushrooms have gills that are thin and white:

False: The data provided strongly disproves this theory. Almost an even split between edible and poisonous mushroom having gills that are thin and white.

Proposed Rules

I could only find one rule not mentioned above to be true. The training data strongly (100%) supports "green spore print color mushroom are poisonous".

Rule 9: Poisonous mushrooms have green spore print color: